

**Solution 1**

---

**Solution 1 (a)****Step 1: Identify characteristics of the dataspace**

We are given 10 dimensional vectors where each element can be any real number ( $x_i \in \mathbb{R}$ ):

$\therefore$  we can express the dataspace  $\chi$  as:  $\chi = \mathbb{R}^{10}$

---

**Solution 1 (b)****Step 1: Identify characteristics of the dataspace**

We are given 3 dimensional vectors where each element is zero or one ( $x_i \in [0, 1]$ ):

$\therefore$  we can express the dataspace  $\chi$  as:  $\chi = [0, 1]^3$

---

## Solution 2

---

### Solution 2 (a)

**Step 1: Define Euclidean distance ( $\ell_2$ )**

$$\ell_2 = \|p - q\|_2 = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

**Step 2: Compute  $\ell_2$**

Let  $p = 1$  and  $q = 10$

$$\begin{aligned}\ell_2 &= \sqrt{\sum_{i=1}^n (p_i - q_i)^2} \\ \ell_2 &= \sqrt{\sum_{i=1}^1 (1 - 10)^2} \\ \ell_2 &= \sqrt{(-9)^2} \\ \ell_2 &= 9\end{aligned}$$

$\therefore \ell_2 = 9$

---

### Solution 2 (b)

**Step 1: Define Euclidean distance ( $\ell_2$ )**

$$\ell_2 = \|p - q\|_2 = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

**Step 2: Compute  $\ell_2$**

Let  $p = \begin{bmatrix} -1 \\ 12 \end{bmatrix}$ ,  $q = \begin{bmatrix} 6 \\ -12 \end{bmatrix}$

$$\begin{aligned}\ell_2 &= \sqrt{\sum_{i=1}^n (p_i - q_i)^2} \\ \ell_2 &= \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2} \\ \ell_2 &= \sqrt{(-1 - 6)^2 + (12 - (-12))^2} \\ \ell_2 &= \sqrt{(-7)^2 + (24)^2} \\ \ell_2 &= \sqrt{625} \\ \ell_2 &= 25\end{aligned}$$

$\therefore \ell_2 = 25$

---

**Solution 2**

---

**Solution 2 (c)****Step 1: Define Euclidean distance ( $\ell_2$ )**

$$\ell_2 = \|p - q\|_2 = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

**Step 2: Compute  $\ell_2$** 

$$\text{Let } p = \begin{bmatrix} 1 \\ 5 \\ -1 \end{bmatrix}, q = \begin{bmatrix} 5 \\ 2 \\ 11 \end{bmatrix}$$

$$\begin{aligned}\ell_2 &= \sqrt{\sum_{i=1}^n (p_i - q_i)^2} \\ \ell_2 &= \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + (p_3 - q_3)^2} \\ \ell_2 &= \sqrt{(1 - 5)^2 + (5 - 2)^2 + (-1 - 11)^2} \\ \ell_2 &= \sqrt{(-4)^2 + (3)^2 + (-12)^2} \\ \ell_2 &= \sqrt{169} \\ \ell_2 &= 13\end{aligned}$$

$$\therefore \ell_2 = 13$$

---

**Solution 3**

---

**Solution 3 (a)****Step 1: Normalize the vector  $x$** 

$$\text{Let } x = \begin{bmatrix} 10 \\ 15 \\ 25 \end{bmatrix}$$

$$\sum_{i=1}^3 x_i = x_1 + x_2 + x_3 = 10 + 15 + 25 = 50$$

Now, divide each entry by the total sum:

$$p = \frac{1}{50} \cdot x = \frac{1}{50} \begin{bmatrix} 10 \\ 15 \\ 25 \end{bmatrix} = \begin{bmatrix} 10/50 \\ 15/50 \\ 25/50 \end{bmatrix} = \begin{bmatrix} 0.2 \\ 0.3 \\ 0.5 \end{bmatrix}$$

$\therefore$  the result ( $p$ ) of scaling vector  $x$  is the following:

$$p = \begin{bmatrix} 0.2 \\ 0.3 \\ 0.5 \end{bmatrix}$$

---

**Solution 3 (b)****Step 1: Define dimension of the probability simplex**

The dimension of vector  $p$  is 3 and  $k = n - 1$  where  $k$  is the dimension of the probability simplex

$\therefore$  vector  $p$  lies in the probability simplex( $\Delta_2$ ) for  $k = 2$

---

**Solution 4**

---

**Step 1: Define probability simplex  $\Delta_2$** 

For a point to be scalable to  $\Delta_2$ , after scaling it must satisfy:

- All components must be non-negative
- The sum of components must equal 1

**Step 2: Give example that violates one of the rules in Step 1**

Let  $x = \begin{bmatrix} 1 \\ -2 \end{bmatrix}$

The second component of  $x$  violates the first rule, all components for a point must be non-negative  $\Delta_2$ .

$\therefore$  the point  $x = \begin{bmatrix} 1 \\ -2 \end{bmatrix}$  cannot be scaled to lie in  $\Delta_2$

---

## Solution 5

---

### Visualizing the Simplex $\Delta_3$ in 2D Projections

Here are the three 2D views of the probability simplex  $\Delta_3$ . Each plot is a *shadow* of the 3D triangle, viewed along one of the principal axes.

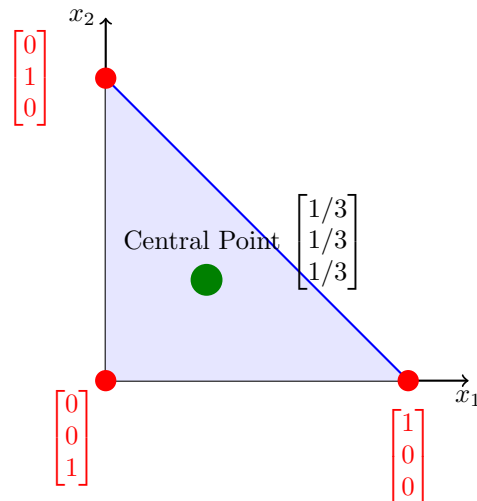


Figure 1: View 1: Projection onto the  $x_1$ - $x_2$  plane.




Figure 2: View 2: Projection onto the  $x_2$ - $x_3$  plane.

**Solution 5**

---




Figure 3: View 3: Projection onto the  $x_1$ - $x_3$  plane.

---

**Solution 6**

---

**6 (a):  $\ell_1$  for  $p$  and  $q$** 

The  $\ell_1$  distance between two vectors  $p, q \in \mathbb{R}^n$  is given by:

$$\|p - q\|_1 = \sum_{i=1}^n |p_i - q_i|$$

$$\text{Let } p = \begin{bmatrix} 1/2 \\ 1/4 \\ 1/8 \\ 1/8 \end{bmatrix} \text{ and } q = \begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix}$$

$$\begin{aligned} \|p - q\|_1 &= \left| \frac{1}{2} - \frac{1}{4} \right| + \left| \frac{1}{4} - \frac{1}{4} \right| + \left| \frac{1}{8} - \frac{1}{4} \right| + \left| \frac{1}{8} - \frac{1}{4} \right| \\ &= \left| \frac{2}{4} - \frac{1}{4} \right| + |0| + \left| \frac{1}{8} - \frac{2}{8} \right| + \left| \frac{1}{8} - \frac{2}{8} \right| \\ &= \frac{1}{4} + 0 + \left| -\frac{1}{8} \right| + \left| -\frac{1}{8} \right| \\ &= \frac{1}{4} + \frac{1}{8} + \frac{1}{8} = \frac{2}{8} + \frac{1}{8} + \frac{1}{8} = \frac{4}{8} = \frac{1}{2} \end{aligned}$$

$$\therefore \ell_1 = \frac{1}{2}$$

---



**Solution 6**

---

**6 (b):  $\ell_1$  for  $q$  and  $r$**

The  $\ell_1$  distance between two vectors  $q, r \in \mathbb{R}^n$  is given by:

$$\|q - r\|_1 = \sum_{i=1}^n |q_i - r_i|$$

$$\text{Let } q = \begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \text{ and } r = \begin{bmatrix} 1/2 \\ 0 \\ 1/4 \\ 1/4 \end{bmatrix}$$

$$\begin{aligned} \|q - r\|_1 &= \sum_{i=1}^4 |q_i - r_i| \\ &= \left| \frac{1}{4} - \frac{1}{2} \right| + \left| \frac{1}{4} - 0 \right| + \left| \frac{1}{4} - \frac{1}{4} \right| + \left| \frac{1}{4} - \frac{1}{4} \right| \\ &= \left| \frac{1}{4} - \frac{2}{4} \right| + \left| \frac{1}{4} \right| + |0| + |0| \\ &= \left| -\frac{1}{4} \right| + \frac{1}{4} + 0 + 0 \\ &= \frac{1}{4} + \frac{1}{4} \\ &= \frac{1}{2} \end{aligned}$$

$$\therefore \ell_1 = \frac{1}{2}$$

---

## Solution 6

---

### 6 (c): KL divergence $K(p, q)$

The Kullback-Leibler (KL) divergence from a distribution  $p$  to a distribution  $q$  is defined as:

$$K(p, q) = \sum_i p_i \ln \frac{p_i}{q_i}$$

$$\text{Let } p = \begin{bmatrix} 1/2 \\ 1/4 \\ 1/8 \\ 1/8 \end{bmatrix} \text{ and } q = \begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix}$$

$$\begin{aligned} K(p, q) &= \sum_{i=1}^4 p_i \ln \left( \frac{p_i}{q_i} \right) \\ &= p_1 \ln \left( \frac{p_1}{q_1} \right) + p_2 \ln \left( \frac{p_2}{q_2} \right) + p_3 \ln \left( \frac{p_3}{q_3} \right) + p_4 \ln \left( \frac{p_4}{q_4} \right) \\ &= \frac{1}{2} \ln \left( \frac{1/2}{1/4} \right) + \frac{1}{4} \ln \left( \frac{1/4}{1/4} \right) + \frac{1}{8} \ln \left( \frac{1/8}{1/4} \right) + \frac{1}{8} \ln \left( \frac{1/8}{1/4} \right) \\ &= \frac{1}{2} \ln(2) + \frac{1}{4} \ln(1) + \frac{1}{8} \ln \left( \frac{1}{2} \right) + \frac{1}{8} \ln \left( \frac{1}{2} \right) \\ &= \frac{1}{2} \ln(2) + \frac{1}{4}(0) - \frac{1}{8} \ln(2) - \frac{1}{8} \ln(2) \\ &= \frac{1}{2} \ln(2) - \frac{2}{8} \ln(2) \\ &= \left( \frac{1}{2} - \frac{1}{4} \right) \ln(2) \\ &= \frac{1}{4} \ln(2) \end{aligned}$$

$$\therefore K(p, q) = \frac{1}{4} \ln(2)$$


---

**Solution 6**

---

**6 (d): KL divergence  $K(q, r)$** 

The Kullback-Leibler (KL) divergence from a distribution  $q$  to a distribution  $r$  is defined as:

$$K(q, r) = \sum_i q_i \ln \frac{q_i}{r_i}$$

$$\text{Let } q = \begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix} \text{ and } r = \begin{bmatrix} 1/2 \\ 0 \\ 1/4 \\ 1/4 \end{bmatrix}$$

Looking at the second component ( $i = 2$ ). Here,  $q_2 = \frac{1}{4} > 0$  while  $r_2 = 0$ . The corresponding term in the KL divergence sum,  $q_2 \ln \left( \frac{q_2}{r_2} \right)$ , involves division by zero.

Hence, the divergence will be infinite.

$$\therefore K(q, r) = \infty$$

---

## Solution 7

---

### Python Code

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from extract_feature import compute_or_load_features
4 from sklearn.neighbors import KNeighborsClassifier
5
6 def run_nearest_neighbor(x_train, y_train, x_test, y_test):
7     # create classifier
8     nn_classifier = KNeighborsClassifier(n_neighbors=1, algorithm='auto')
9
10    # train
11    nn_classifier.fit(x_train, y_train)
12
13    # test and report accuracy
14    test_acc = nn_classifier.score(x_test, y_test)
15
16    print("Nearest neighbor accuracy on the test set: %f"%test_acc)
17
18    return nn_classifier
19
20
21 def analyze_nn(classifier, x_train, y_train, x_test, y_test, x_test_features, N,
22               model_type):
23     """
24     generalization to grab indices, predictions, and make plots
25     """
26     # list of image labels
27     CIFAR_CLASSES = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog',
28                      'horse', 'ship', 'truck']
29
30     # get predictions of classifier on test data
31     y_pred = classifier.predict(x_test_features)
32
33     # define bool condition where classifier made correct prediction
34     is_correct = (y_pred == y_test)
35
36     # get indices for correct and incorrect samples
37     correct_indices_test = np.where(is_correct)[0][:N]
38     incorrect_indices_test = np.where(~is_correct)[0][:N]
39
40     # make image plots
41     def plot_pairs(title, test_indices):
42         # check just in case of bad result.. nema nista
43         if len(test_indices) == 0:
44             return
45
46         # Get the feature vectors for the selected test images (flattened)
47         selected_test_features = x_test_features[test_indices]
48
49         # get nearest neighbor
50         # note: kneighbors -> (distances, indices).. only need the indices.
51         _, nn_train_indices_2D = classifier.kneighbors(X=selected_test_features,
52                                                       n_neighbors=1)
53
54         # 2D -> 1D
55         nn_train_indices = nn_train_indices_2D.flatten()
56
57         # get images and labels for the selected test points
58         test_images_sample = x_test[test_indices]
59         test_labels_sample = y_test[test_indices]
60
61         # get the RAW neighbor image and label from the training points
62         nn_images_sample = x_train[nn_train_indices]
63         nn_labels_sample = y_train[nn_train_indices]

```

```

63     # prediction is nearest neighbor label
64     y_pred_sample = nn_labels_sample
65
66     # reshape (N, C, H, W) to (N, H, W, C) for plot
67     test_images_plt = test_images_sample.transpose(0, 2, 3, 1)
68     nn_images_plt = nn_images_sample.transpose(0, 2, 3, 1)
69
70     N_plot = len(test_indices)
71
72     # just in case only 1 image is plotted (axes will be 1D instead of 2D)
73     if N_plot == 1:
74         fig, axes = plt.subplots(N_plot, 2, figsize=(6, 2))
75         axes = axes[np.newaxis, :]
76     else:
77         fig, axes = plt.subplots(N_plot, 2, figsize=(6, 2 * N_plot))
78
79     fig.suptitle(title, fontsize=14, y=1.02)
80
81     for i in range(N_plot):
82         is_correct = (test_labels_sample[i] == y_pred_sample[i])
83         pred_color = 'green' if is_correct else 'red'
84
85         # plot test image
86         axes[i, 0].imshow(test_images_plt[i] / 255.0)
87         axes[i, 0].set_title(f"Test ({CIFAR_CLASSES[test_labels_sample[i]]})",
88                             fontsize=10)
89         axes[i, 0].axis('off')
90
91         # plot nearest neighbor
92         axes[i, 1].imshow(nn_images_plt[i] / 255.0)
93         axes[i, 1].set_title(
94             f"NN ({CIFAR_CLASSES[nn_labels_sample[i]]})",
95             fontsize=10,
96             color=pred_color
97         )
98         axes[i, 1].axis('off')
99
100     plt.tight_layout(rect=[0, 0.03, 1, 0.98])
101     plt.show()
102     fig.savefig(f"{title.split(' ')[2].lower()}_{model_type}.png")
103
104     # plot correct and incorrect cases
105     plot_pairs(f"First {N} Correct Predictions ({model_type})", correct_indices_test)
106     plot_pairs(f"First {N} Incorrect Predictions ({model_type})",
107               incorrect_indices_test)
108
109 # raw pixel
110 raw_pixel_train_features, raw_pixel_test_features = compute_or_load_features(x_train,
111                                     x_test, "raw_pixel")
112 raw_pixel_knn_classifier = run_nearest_neighbor(raw_pixel_train_features, y_train,
113         raw_pixel_test_features, y_test)
114 analyze_nn(
115     classifier=raw_pixel_knn_classifier,
116     x_train=x_train,
117     y_train=y_train,
118     x_test=x_test,
119     y_test=y_test,
120     x_test_features=raw_pixel_test_features,
121     N=5,
122     model_type="raw_pixel"
123 )
124
125 # HoG
126 hog_train_features, hog_test_features = compute_or_load_features(x_train, x_test, "hog")
127 hog_knn_classifier = run_nearest_neighbor(hog_train_features, y_train,
128         hog_test_features, y_test)
129 analyze_nn(
130     classifier=hog_knn_classifier,
131     x_train=x_train,
132     y_train=y_train,
133     x_test=x_test,
134     y_test=y_test,
135     x_test_features=hog_test_features,

```

```
131     N=5,
132     model_type="hog"
133 )
134
135 # vgg-last-fc
136 pretrained_cnn_last_fc_train_features, pretrained_cnn_last_fc_test_features =
137     compute_or_load_features(x_train, x_test, "pretrained_cnn", "last_fc")
138 pretrained_cnn_last_fc_knn_classifier =
139     run_nearest_neighbor(pretrained_cnn_last_fc_train_features, y_train,
140         pretrained_cnn_last_fc_test_features, y_test)
141 analyze_nearest_neighbors_simple(
142     classifier=pretrained_cnn_last_fc_knn_classifier,
143     x_train=x_train,
144     y_train=y_train,
145     x_test=x_test,
146     y_test=y_test,
147     x_test_features=pretrained_cnn_last_fc_test_features,
148     N=5,
149     model_type='vgg-last-fc'
150 )
```

**Solution 7**

**Part a: Dimensionality for each of the representations (raw pixel, HoG, VGG-last-fc, VGG-last-conv)**

Feature Type	Dimensionality
Raw Pixel	3072
HoG	512
VGG-last-fc	4096
VGG-last-conv	512

---

**Part b: Test accuracies for 1-nearest neighbor classification using the various representations (raw pixel, HoG, VGG-last-fc, VGG-last-conv, random-VGG-last-fc, random-VGG-last-conv).**

Feature Type	1-NN test accuracy (%)
Raw Pixel	35.4
HoG	36.6
VGG-last-fc	92.1
VGG-last-conv	92.0
random VGG-last-fc	39.1
random VGG-last-conv	40.6

---

## Solution 7

### Part c: Raw Pixel correct/incorrect




Figure 4: First five correct/incorrect images for Raw Pixel



## Solution 7

### Part c: HoG correct/incorrect




Figure 5: First five correct/incorrect images for HoG

## Solution 7

### Part c: VGG-last-fc correct/incorrect

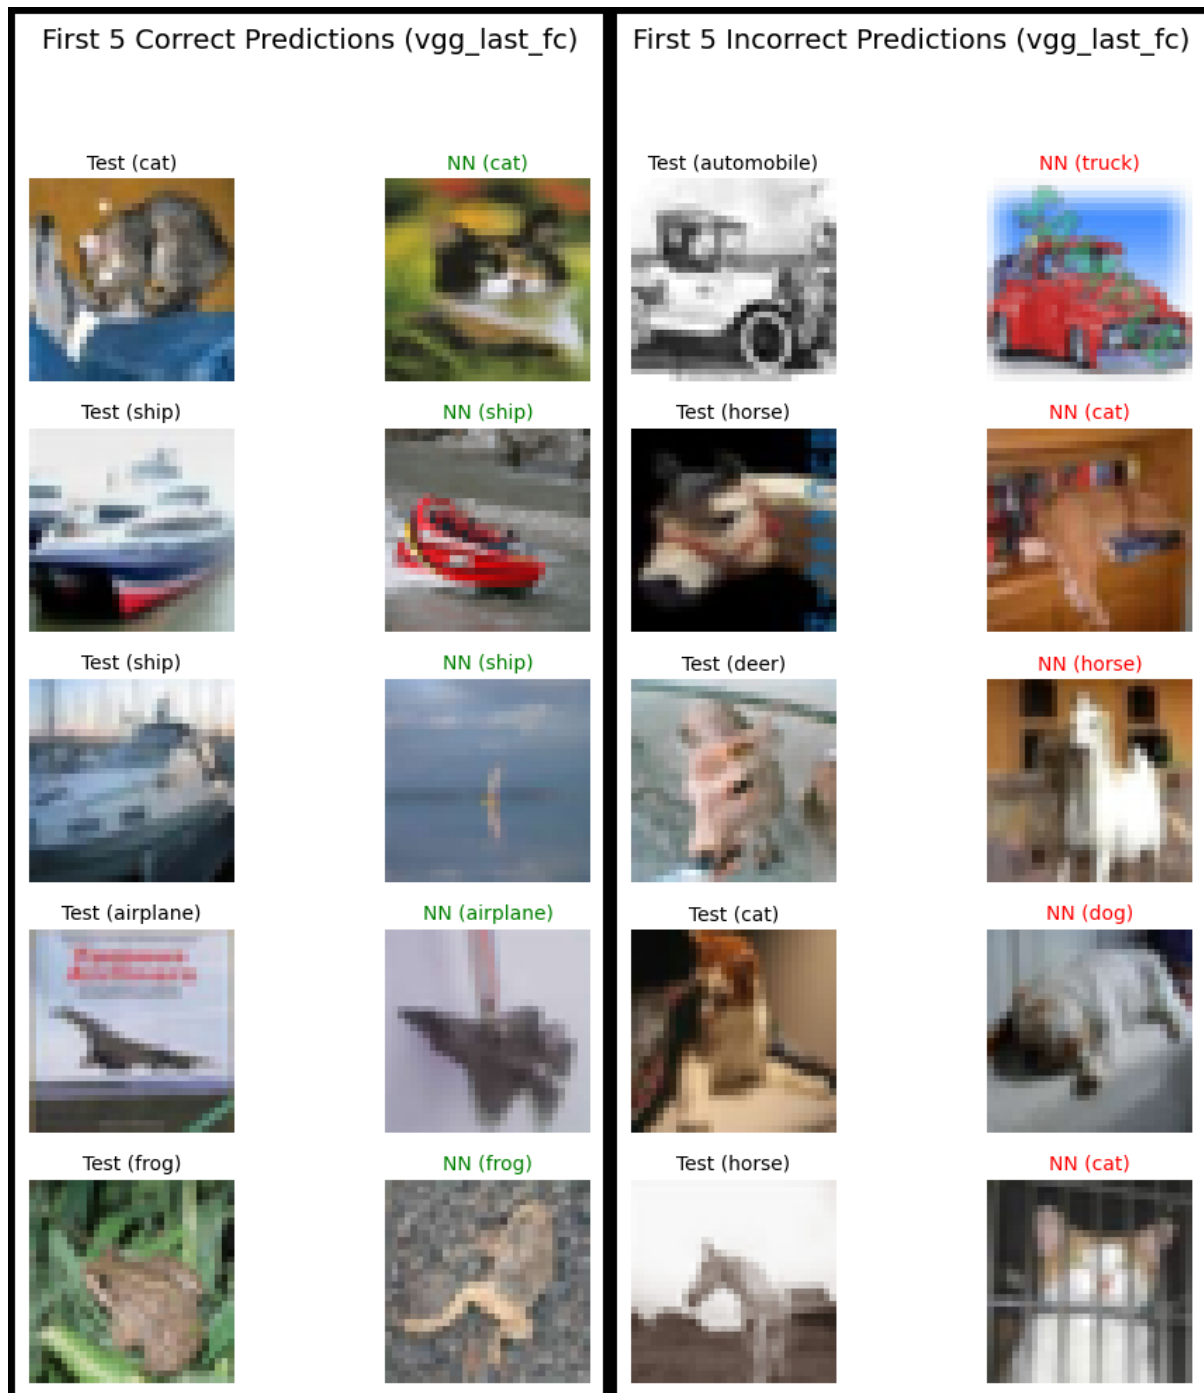


Figure 6: First five correct/incorrect images for VGG-last-fc

## Solution 8

---

### Python Code

```
1 import numpy as np
2 from sklearn.neighbors import NearestNeighbors
3
4 filename = 'glove.6B.300d.txt'
5 with open(filename) as f:
6     content = f.read().splitlines()
7
8 # initialize vecs and words 'containers'
9 n = len(content)
10 vecs = np.zeros((n, 300))
11 words = [" " for i in range(n)]
12 for index, rawline in enumerate(content):
13     line = rawline.split()
14     words[index] = line[0]
15     # need numpy speed
16     vecs[index] = np.array(line[1:], dtype=np.float32)
17
18 # make dict for access to word and index
19 word_to_index = {word: i for i, word in enumerate(words)}
20
21 # initialize target words
22 target_words = ['communism', 'africa', 'happy', 'sad', 'upset', 'computer', 'cat',
23                 'dollar']
24
25 # find 5 nearest neighbors (n).. remember k = n+1
26 n_neighbors = 6
27
28 # initialize and fit the nn model
29 nn_model = NearestNeighbors(n_neighbors=n_neighbors, metric='euclidean',
30                             algorithm='auto')
31 nn_model.fit(vecs)
32
33 # gracefully check for typo
34 try:
35     target_indices = [word_to_index[word] for word in target_words]
36 except KeyError as e:
37     print(f" error: word '{e.args[0]}' was not found.")
38     exit()
39
40 # extract the corresponding vectors for the target words from vecs
41 target_vectors = vecs[target_indices]
42
43 # find the nearest neighbors
44 distances, indices = nn_model.kneighbors(target_vectors)
45
46 # format and print results
47 results = {}
48 for i, word in enumerate(target_words):
49     neighbor_indices = indices[i][1:]
50     neighbor_words = [words[idx] for idx in neighbor_indices]
51     results[word] = neighbor_words
52
53 print(f"5 nearest neighbors in {filename} for {target_words}")
54 print(results)
```

## Solution 8

---

### Word Vectors: 5 closest words

Target Word	Five Closest Words
communism	['fascism', 'capitalism', 'nazism', 'stalinism', 'socialism']
africa	['african', 'continent', 'south', 'africans', 'zimbabwe']
happy	['glad', 'pleased', 'always', 'everyone', 'sure']
sad	['sorry', 'tragic', 'happy', 'pathetic', 'awful']
upset	['upsetting', 'surprised', 'upsets', 'stunned', 'shocked']
computer	['computers', 'software', 'technology', 'laptop', 'computing']
cat	['cats', 'dog', 'pet', 'feline', 'dogs']
dollar	['currency', 'dollars', 'euro', 'multibillion', 'weaker']

---