

ONLINE MASTERS IN DATA SCIENCE

DSC 257R - UNSUPERVISED LEARNING

# *K*-MEANS CLUSTERING

SANJOY DASGUPTA, PROFESSOR


UC San Diego

COMPUTER SCIENCE & ENGINEERING  
HALICIOĞLU DATA SCIENCE INSTITUTE

## The $K$ -Means Optimization Problem

- Input: Points  $x_1, \dots, x_n \in \mathbb{R}^d$ ; integer  $k$
- Output: "Centers", or representatives,  $\mu_1, \dots, \mu_k \in \mathbb{R}^d$
- Goal: Minimize average squared distance between points and their nearest representatives:


$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



## The $K$ -Means Optimization Problem

- Input: Points  $x_1, \dots, x_n \in \mathbb{R}^d$ ; integer  $k$
- Output: "Centers", or representatives,  $\mu_1, \dots, \mu_k \in \mathbb{R}^d$
- Goal: Minimize average squared distance between points and their nearest representatives:


$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



## The $K$ -Means Optimization Problem



- Input: Points  $x_1, \dots, x_n \in \mathbb{R}^d$ ; integer  $k$
- Output: "Centers", or representatives,  $\mu_1, \dots, \mu_k \in \mathbb{R}^d$
- Goal: Minimize average squared distance between points and their nearest representatives:

$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



Centers carve  $\mathbb{R}^d$  into  $k$  **convex** regions:  
 $\mu_j$ 's region consists of points for which  
it is the closest center.

## An Unfavorable Optimization Landscape




In fact,  $k$ -means is an **NP-hard** optimization problem.

What can we hope for in such situations?


## Heuristic: Lloyd's $K$ -Means Algorithm

- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.




## Heuristic: Lloyd's $K$ -Means Algorithm

- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



## Heuristic: Lloyd's $K$ -Means Algorithm


- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.






## Heuristic: Lloyd's $K$ -Means Algorithm

- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.




## Heuristic: Lloyd's $K$ -Means Algorithm

- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.




## Heuristic: Lloyd's $K$ -Means Algorithm

- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.





## Heuristic: Lloyd's $K$ -Means Algorithm

- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



Each iteration reduces the cost  $\Rightarrow$  convergence to a local optimum.

## Initialization Matters



## Initializing the $K$ -Means Algorithm

Typical practice: choose  $k$  data points at random as the initial centers.

## Initializing the $K$ -Means Algorithm

Typical practice: choose  $k$  data points at random as the initial centers.

Another common trick: start with extra centers, then prune later.

## Initializing the $K$ -Means Algorithm

Typical practice: choose  $k$  data points at random as the initial centers.

Another common trick: start with extra centers, then prune later.

A particularly good initializer:  **$k$ -means++**

- Pick a data point  $x$  at random as the first center
- Let  $C = \{x\}$  (centers chosen so far)
- Repeat until desired number of centers is attained:
  - Pick a data point  $x$  at random from the following distribution:

$$\Pr(x) \propto \text{dist}(x, C)^2,$$

where  $\text{dist}(x, C) = \min_{z \in C} \|x - z\|$

- Add  $x$  to  $C$



